

The development of abstract concepts in the early lexical network

Anonymous NAACL submission

Abstract

How do children learn abstract concepts such as animal vs. artefact? Previous research suggests that such concepts can be derived using cues from the language they hear around them, e.g., word co-occurrence. We propose that such information can be integrated in the evolving lexical network, and higher-level organization can be identified bottom-up as the dense components in this network. We found that early abstract categories develop simultaneously thanks to the children's word trajectory which favors the exploration of the global conceptual space.

1 Introduction:

One of the central challenges in cognitive development is to understand how concepts develop (Carey, 2009; Keil, 1992; Gopnik and Meltzoff, 1997). Of particular interest is the case of abstract concepts which have non-obvious shared properties such as "animal" and "artefacts". For example, a cat and a bird are perceptually quite different but they share some fundamental properties (e.g., breathing, feeding, and reproducing) which make them animals, and which set them apart from the category of artefacts. In such cases learning requires, at least in part, cultural/linguistic cues which provide information beyond what can be obtained through the senses (Gelman, 2009; Harris, 2012; Csibra and Gergely, 2009).

One way children's conceptual learning can benefit from the language they hear around them is through word co-occurrence. For example, one can learn an abstract concept (e.g., animal) simply by observing how its instances (e.g., "cat" and "bird") go together in speech. This is a plausible learning strategy. On the one hand, experi-

mental research has shown that children track co-occurrence statistics (Saffran et al., 1996). On the other hand, inspection of the caregiver's input shows that it contains rich co-occurrence information about various abstract concepts (Huebner and Willits, 2018).

What remains to be investigated is the way abstract concepts develop from the interaction of the children's learning mechanisms and the structure of their linguistic input. We study this development and compare it to two hypothetical developmental scenarios. On the first, learning starts by exploring the global conceptual structure; categories are refined simultaneously over development (we call this mechanism exploration-based learning). On the second scenario, learning starts by exploring a small region of the conceptual space (e.g., the category "animals") and only after the refinement of this category, does the learner move to another (we call this mechanism exploitation-based learning).

The paper is organized as follows. First we describe the research strategy. In brief, we represented the developing lexicon as an evolving network and we used word co-occurrence in parent speech as a measure of words' relatedness. We operationalized high-level concepts as the highly interconnected components of the network. Second, we explore how the pattern of children's word learning influences higher-level conceptual development, and we compare this pattern to the development induced by the exploration-based and exploitation-based mechanisms. Finally, we discuss the findings in light of previous research.

2 Data and Methods

2.1 Constructing lexical networks

The networks' nodes were nouns from Wordbank (Frank et al., 2017), an open repository aggregat-

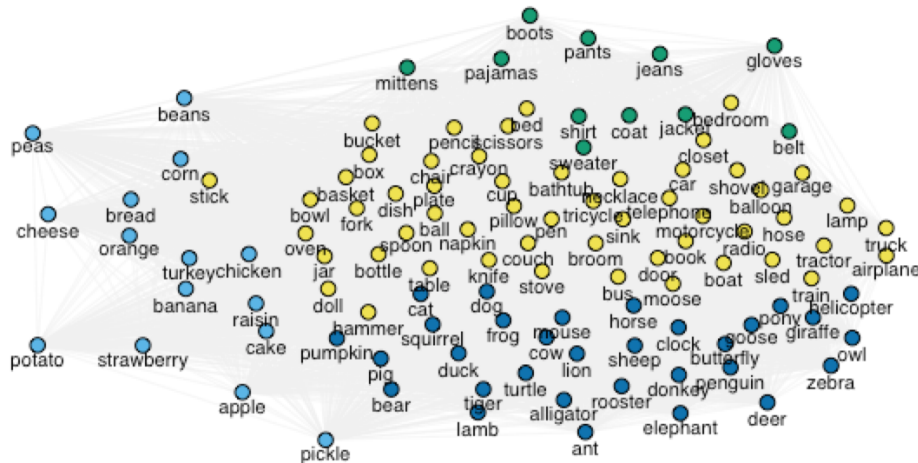


Figure 1: Network obtained using a sample of nouns in CDI data (nodes), and co-occurrence-based similarity from a corpus of child-directed speech (edges). Colors indicate highly interconnected clusters identified using unsupervised network community detection. The clusters correspond, overall, to four higher-level concepts: animal, food, clothes, and artefacts.

ing cross-linguistic language developmental data of the MacArthur-Bates Communicative Development Inventory (CDI), a parent report vocabulary checklist, Toddler version (Fenson et al., 1994). Pairs of nouns were linked by weighted edges representing their semantic similarity derived based on co-occurrence in the corpus of child directed speech CHILDES (MacWhinney, 2014), using Word2Vec algorithm (Mikolov et al., 2013).

First, we constructed the end-state network based on a subset of CDI data. This subset was made of nouns belonging to a diversity of semantic categories. Items of this subset (named “uni_lemmas” in the WordBank database) were translated across the languages used in this study, allowing us to account for cross-linguistic variability. Second, to study development towards the end-state, we constructed a different network at each month, based on the nouns that have been learned by that month.

2.2 Identifying high-level concepts in a network

We assume that high-level concepts correspond to clusters of highly interconnected nodes in the networks. We identified such clusters using WalkTrap (Pons and Latapy, 2006), an unsupervised community detection algorithm based on the fact that a random walker tends to be trapped in dense parts of a network. Figure 1 shows the outcome of clus-

ter identification in the end-state network. The algorithm obtained four major clusters corresponding to the categories of clothes, food, animal and artefacts. We refer to this end-state clustering as C^* . To examine developmental change in the conceptual organization, we run the cluster identification algorithm at each month of acquisition t , and we compare the resulting clustering, noted C_t , to that of the end-state C^* . The method of this comparison is detailed below.

2.3 Measuring conceptual development

We measure conceptual development by comparing \mathcal{C}_t to \mathcal{C}^* across time. We used a standard method in clustering comparison, which is based on counting word pairs on which the two clusterings agree or disagree (Rand, 1971; Hubert and Arabie, 1985). A pair of words learned by month t can fall under one of the four following cases:

1. True positives $tp(\mathcal{C}_t)$: pairs that are placed in the same cluster under \mathcal{C}_t and in the same cluster under \mathcal{C}^* .
2. True negatives $tn(\mathcal{C}_t)$: pairs placed in different clusters under \mathcal{C}_t and in different clusters under \mathcal{C}^* .
3. False positive $fp(\mathcal{C}_t)$: pairs placed in the same cluster under \mathcal{C}_t and in different clusters under \mathcal{C}^* .




Figure 2: Mean precision and recall scores obtained through comparing the end-state clustering to clusterings at different months of acquisition, across different languages. Colors indicates real and hypothetical word sampling mechanisms. Errors bars represent 95% confidence intervals.

4. False negatives $fn(\mathcal{C}_t)$: pairs placed in different clusters under \mathcal{C}_t and in the same cluster under \mathcal{C}^* .

We quantify clustering comparison using precision $P(\mathcal{C}_t)$ and recall $R(\mathcal{C}_t)$, defined as follows:

$$P(\mathcal{C}_t) = \frac{|tp(\mathcal{C}_t)|}{|tp(\mathcal{C}_t)| + |fp(\mathcal{C}_t)|}$$

$$R(\mathcal{C}_t) = \frac{|tp(\mathcal{C}_t)|}{|tp(\mathcal{C}_t)| + |fn(\mathcal{C}_t)|}$$

We made this comparison using different degrees of clustering granularity. More precisely, we fixed the same number of clusters for both \mathcal{C}_t and \mathcal{C}^* , and we varied this number from two to four clusters. We did not use the trivial case of one cluster, nor did we use more than four clusters, since this number was optimal for the largest network (i.e., the end-state network) based on the modularity maximization criterion (Newman, 2006).

2.4 Learning mechanisms

We examined how higher-order concepts develop under the children’s real word learning trajectory. We compared this development to the development induced by an exploration- and an exploitation-like learning. To construct the real word learning trajectory, we used the normative age of acquisition, that is, the age at which a word is produced by at least 50% of children in each language (Goodman et al., 2008).

The exploration-based learning was instantiated as a uniform sampling from the end-state vocabulary. The exploitation-based learning had the additional constraint of sampling from one category at a time: the first word is selected randomly from one cluster, subsequent words are sampled from the same cluster. After all words from this cluster are used, a word from a different cluster is chosen, and the same process is repeated until all clusters are covered.¹

3 Results

Figure 2 shows the scores obtained through comparing \mathcal{C}^* to \mathcal{C}_t at different points in time t . For the real word learning trajectory, both precision and recall start relatively low, indicating that the induced conceptual organization is initially quite different from that of the end-state. Both measures converge towards 1 (i.e., perfect score) as \mathcal{C}_t becomes more and more similar to \mathcal{C}^* .

We first compared the real conceptual development to the one induced by exploration-based learning. To this end, we fit a mixed-effect regression using age and learning (real vs. exploration-based) as fixed effects, and using language and number of fixed clusters as random effects accounting for the nested structure of the data. Real

¹Note that the way we instantiated the exploitation-based learning is not fully unsupervised, but we were more interested in modeling extreme cases to which real learning can be compared.

and exploration-based patterns of conceptual developmental were indistinguishable for both precision ($\beta = 0.12$, $p = 0.09$) and recall ($\beta = 0.09$, $p = 0.1$).

Similarly, we compared real conceptual development to exploitation-based learning. Using a similar regression, we found a difference in both precision ($\beta = 1.08$, $p < 0.01$) and recall ($\beta = -0.2$, $p < 0.01$). Exploitation-base learning had generally higher precision, indicating there to be less false positive pairs. This result is due to the fact that we sampled instances from a same category. However, the same type of learning had lower recall, indicating there to be more false negative pairs. This second result was due to the fact that sampling from a same category leads to clusterings that are finer in their conceptual granularity than the end-state.

4 Discussion

This paper asks whether children can learn abstract concepts based on word co-occurrence in the language they hear around them. We found that, when using co-occurrence information in the developing lexical network, several high-level concepts such as “animal”, “artefact”, “food” and “clothes” emerge bottom-up as clusters of highly inter-connected nodes. In addition, these categories develop simultaneously thanks to the children’s word learning trajectory which tends to favor the exploration of the global conceptual landscape rather than the exploitation and refinement of one specific category at time.

The development of the higher-level conceptual structure seems to be unaffected by the order with which words are acquired (as long as this order approximates a uniform/exploration-like sampling), suggesting that the process of conceptual development can accommodate a wide range of word learning trajectories without qualitative change in the higher order organization. For example, whether acquisition start first with the words “cat” and “banana” or with the words “cow” and “potato” does not qualitatively affect the higher-level organization involving “animal” and “food”. This property is important as it suggests, for instance, that development may be resilient to variability in the children’s linguistic input (Slobin, 2014; Hart and Risley, 1995).

Developmental changes were captured by precision and recall. The increase in precision means

that false positives decrease over time: some word pairs that are initially lumped together in a same category, are eventually differentiated. Similarly, the increase in recall means that false negatives decrease, that is, some word pairs that are initially distinct, become eventually subsumed by a same category. These patterns suggest a process of conceptual reorganization involving both “differentiation” and “coalescence” as was suggested in the developmental literature (Carey, 2009).

That said, these developmental changes were not necessarily related to specific concepts (since the patterns were similar when we randomized the order of word learning). Instead, this finding suggests that differentiation and coalescence of word pairs in our data are related to the change in the vocabulary size across development: As more words are added to their lexical network, learners may approximate better the underlying conceptual organization of the mature lexicon, and would make less categorization errors. Indeed, research in network science indicates that properties of a real networks become more distorted as the size of a sampled sub-network decreases (Leskovec and Faloutsos, 2006).

One limitation of this study is that we used the normative age of acquisition, computed using different children at different age groups. This choice was due to the cross-sectional nature of available CDI data. Though such a measure has been widely used to study important aspects of the early lexical network (Hills et al., 2009; Stella et al., 2017; Storkel, 2009), it only applies at the population level. In our case, though we found that average concept development is driven more by an exploration-based mechanism, some individual children may display an exploitation-like behavior. For example, prior knowledge about dinosaurs may enable the learning of new dinosaur-related words more easily (Chi and Koeske, 1983).

In sum, this work provided a quantitative account of how abstract concepts can emerge from the interaction of the children’s learning mechanism and the properties of their linguistic input. One important direction for future work is to investigate the extent to which the correlational findings obtained in this study (e.g., the identity of categories formed across development or the fact that categorization errors decrease with the size of the lexicon) can be corroborated by controlled behavioral experiments.

References

- Susan Carey. 2009. *The origin of concepts*. Oxford University Press.
- Micheline TH Chi and Randi Daimon Koeske. 1983. Network representation of a child’s dinosaur knowledge. *Developmental psychology*, 19(1).
- Gergely Csibra and György Gergely. 2009. Natural pedagogy. *Trends in cognitive sciences*, 13(4).
- Larry Fenson, Philip S. Dale, J. Steven Reznick, Elizabeth Bates, Donna J. Thal, Stephen J. Pethick, Michael Tomasello, Carolyn B. Mervis, and Joan Stiles. 1994. Variability in early communicative development. *Monographs of the Society for Research in Child Development*, 59(5).
- Michael C. Frank, Mika Braginsky, Daniel Yurovsky, and Virginia A. Marchman. 2017. Wordbank: an open repository for developmental vocabulary data. *Journal of Child Language*, 44(3):677–694.
- Susan A Gelman. 2009. Learning from others: Children’s construction of concepts. *Annual review of psychology*, 60.
- Judith C. Goodman, Philip S. Dale, and Ping Li. 2008. Does frequency count? parental input and the acquisition of vocabulary. *Journal of Child Language*, 35(3):515–531.
- Alison Gopnik and Andrew N Meltzoff. 1997. *Words, thoughts, and theories*. MIT Press.
- Paul L Harris. 2012. *Trusting what you’re told: How children learn from others*. Harvard University Press.
- Betty Hart and Todd R Risley. 1995. *Meaningful differences in the everyday experience of young American children*. Paul H Brookes Publishing.
- Thomas T. Hills, Mounir Maouene, Josita Maouene, Adam Sheya, and Linda Smith. 2009. Longitudinal analysis of early semantic networks: Preferential attachment or preferential acquisition? *Psychological Science*, 20(6):729–739.
- Lawrence Hubert and Phipps Arabie. 1985. Comparing partitions. *Journal of classification*, 2(1).
- Philip A Huebner and Jon A Willits. 2018. Structured semantic knowledge can emerge automatically from predicting word sequences in child-directed speech. *Frontiers in Psychology*, 9:133.
- Frank C Keil. 1992. *Concepts, kinds, and cognitive development*. MIT Press.
- Jure Leskovec and Christos Faloutsos. 2006. Sampling from large graphs. In *Proceedings of the 12th ACM SIGKDD international conference on Knowledge discovery and data mining*. ACM.
- Brian MacWhinney. 2014. *The CHILDES project: Tools for analyzing talk, Volume II*. Psychology Press.
- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. 2013. Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems*.
- Mark EJ Newman. 2006. Modularity and community structure in networks. *Proceedings of the national academy of sciences*, 103(23).
- Pascal Pons and Matthieu Latapy. 2006. Computing communities in large networks using random walks. *Journal of Graph Algorithms and Applications*, 10(2).
- William M Rand. 1971. Objective criteria for the evaluation of clustering methods. *Journal of the American Statistical association*, 66(336).
- Jenny R Saffran, Richard N Aslin, and Elissa L Newport. 1996. Statistical learning by 8-month-old infants. *Science*, 274(5294):1926–1928.
- Dan Isaac Slobin. 2014. *The crosslinguistic study of language acquisition*, volume 4. Psychology Press.
- Massimo Stella, Nicole M Beckage, and Markus Brede. 2017. Multiplex lexical networks reveal patterns in early word acquisition in children. *Scientific Reports*, 7.
- Holly L. Storkel. 2009. Developmental differences in the effects of phonological, lexical and semantic variables on word learning by infants. *Journal of Child Language*, 36(2):29–321.